

USER PROMPT:
What's Brazil's capital?

LLM RESPONSE:
Brazil's capital is **Brasília**.

Model: gemini-2.5-flash-lite

Token usage (estimated):

Input tokens: 6
Output tokens: 8
Total tokens: 14

USER PROMPT:
What's Brazil's capital?

LLM RESPONSE:
Brazil's capital is **Brasília**.

Model: gemini-2.5-flash

Token usage (estimated):

Input tokens: 6
Output tokens: 8
Total tokens: 14

USER PROMPT:
What's Brazil's capital?

LLM RESPONSE:
The capital of Brazil is **Brasília**.

Model: gemini-3-flash-preview

Token usage (estimated):

Input tokens: 6
Output tokens: 9
Total tokens: 15

=====

USER PROMPT:

Find the user intent in the following text:
I'm looking for a restaurant around São Paulo who has a good rating for Japanese food.

LLM RESPONSE:
The user intent is to **find a highly-rated Japanese restaurant in São Paulo**.

Model: gemini-2.5-flash-lite

Token usage (estimated):

Input tokens: 33
Output tokens: 19
Total tokens: 52

USER PROMPT:

Find the user intent in the following text:
I'm looking for a restaurant around São Paulo who has a good rating for Japanese food.

LLM RESPONSE:
The user intent is to **find a highly-rated Japanese restaurant in São Paulo**.

Model: gemini-2.5-flash

Token usage (estimated):

Input tokens: 33
Output tokens: 19
Total tokens: 52

USER PROMPT:

Find the user intent in the following text:
I'm looking for a restaurant around São Paulo who has a good rating for Japanese food.

LLM RESPONSE:
The user intent is **to find recommendations for highly-rated Japanese restaurants in São Paulo**.

Specifically, the intent can be broken down into:
* **Goal:** Seek a recommendation/Search for a business.
* **Category:** Food & Drink (Restaurants).

* **Cuisine:** Japanese.
* **Location:** São Paulo.
* **Criteria:** High quality/Positive reviews (good ratings).

Model: gemini-3-flash-preview

Token usage (estimated):

Input tokens: 33
Output tokens: 93
Total tokens: 126

=====

USER PROMPT:

What's Brazil's capital? Respond only with the city name.

LLM RESPONSE:

Brasília

Model: gemini-2.5-flash-lite

Token usage (estimated):

Input tokens: 14
Output tokens: 2
Total tokens: 16

USER PROMPT:

What's Brazil's capital? Respond only with the city name.

LLM RESPONSE:

Brasília

Model: gemini-2.5-flash

Token usage (estimated):

Input tokens: 14
Output tokens: 2
Total tokens: 16

USER PROMPT:

What's Brazil's capital? Respond only with the city name.

LLM RESPONSE:

Brasília

Model: gemini-3-flash-preview

Token usage (estimated):

Input tokens: 14
Output tokens: 2
Total tokens: 16
